

Distillation and autoregression in video generation

Ilya Drobyshevskiy
Yandex Research



Why we want to distill T2V models

Even a short video produce very long token sequence

- A 5-second, 16 FPS video at 480x832 already becomes 33k tokens
- Increasing the resolution to 720x1248 raises this to 73k
- Generating the first video with 1.3B model takes 140s

Methods of the distillation T2V models

VideoLCM: Video Latent Consistency Model

Xiang Wang^{1*} Shiwei Zhang^{2†} Han Zhang³ Yu Liu² Yingya Zhang² Changxin Gao¹ Nong Sang^{1†}

¹Key Laboratory of Image Processing and Intelligent Control,

School of Artificial Intelligence and Automation, Huazhong University of Science and Technology

²Alibaba Group ³Shanghai Jiao Tong University

DOLLAR: Few-Step Video Generation via Distillation and Latent Reward Optimization

Zihan Ding^{1†}, Chi Jin¹, Difan Liu², Haitian Zheng², Krishna Kumar Singh²,

Qiang Zhang², Yan Kang², Zhe Lin², Yuchen Liu^{2†}

¹Princeton University, ²Adobe Research

¹{zihand, chi}@princeton.edu

²{diliu, hazheng, krishsin, qiangz, yankang, zlin, yuliu}@adobe.com

miracle.github.io/dollar/

Dual-Expert Consistency Model for Efficient and High-Q

Zhengyao Lv^{2,3*} Chenyang Si^{1‡*} Tianlin Pan^{1,4}

Kwan-Yee K. Wong² Yu Qiao³ Ziwei Li

¹Nanjing University ²The University of Hong Kong ³Shanghai Arti

⁴University of Chinese Academy of Sciences ⁵S-Lab, Nanyang T

SF-V: Single Forward Video Generation Model

Zhixing Zhang^{1,2*} Yanyu Li¹ Yushu Wu¹ Yanwu Xu¹ Anil Kag¹
Ivan Skorokhodov¹ Willi Menapace¹ Aliaksandr Siarohin¹ Junli Cao¹

Dimitris Metaxas² Sergey Tulyakov¹ Jian Ren^{1†}

¹Snap Inc. ²Rutgers University

Project Page: <https://snap-research.github.io/SF-V>

Motion Consistency Model: Accelerating Video Diffusion with Angled Motion-Appearance Distillation

^{†*} Kevin Lin[†] Zhengyuan Yang[†] Linjie Li[†] Jianfeng Wang[†]

ing Lin[†] David Doermann[†] Junsong Yuan[†] Lijuan Wang[†]

[†]State University of New York at Buffalo [‡]Microsoft

{yzhai6, doermann, jsyuan}@buffalo.edu

{keli, zhengyang, lindsey.li, jianfeng, chungching.lin, lijuanw}@microsoft.com

<https://yhzhai.github.io/mcm/>

TOWARDS ONE-STEP CAUSAL VIDEO ADVERSARIAL SELF-DISTILLATION

Yongqi Yang^{1*} Huayang Huang^{1,*,} Xu Peng^{2,*,} Xiaobin Hu[†]

Jiangning Zhang², Chengjie Wang², Yu Wu^{3†}

¹School of Computer Science, Wuhan University ²Tencent Y

³School of Artificial Intelligence, Wuhan University

{yongqiyang, hyhuang, wuyucs}@whu.edu.cn

Project page: <https://github.com/BigAandSmallq/>

Accelerating Video Diffusion Models via Distribution Matching

Yuanzhi Zhu^{1‡} Hanshu Yan^{1†} Huan Yang¹ Kai Zhang² Junnan Li^{1‡}


¹Rhymes.AI, ²Nanjing University

Adversarial Post-Training for One-Step Video Generation

hanchuan Lin[†] Xin Xia[†] Yuxi Ren[†] Ceyuang Yang[†] Xuefeng Xiao[†] Lu Jiang[†]


DMD/DMD2 recap

Let's combine GAN and DMD losses



CausVid: DMD2 as is

Just swap image with video



How to generate long videos?

- Fast generation of 5-second videos is already possible
- The next challenge is scaling to longer videos
- In practice we often do not know the desired duration beforehand

There are two main approaches:


Key frames interpolation
(less commonly used in prior works)

Autoregressive frames generation
(dominant direction)

Key frames interpolation



Autoregressive frames generation | Teacher Forcing




(a) Teacher Forcing Training


Problems:

- 1) Train-inference mismatch
- 2) Error accumulation

Autoregressive frames generation | Diffusion Forcing



(a) Teacher Forcing Training




(b) Diffusion Forcing Training


DF is trying to fix the train inference mismatch by conditioning not on clean but on noise frames

Diffusion Forcing Sampling

Mixed-noise training allows different inference schedules across frames or chunks. Speedups depend on the sampling schedule and reuse of cached computation.



Autoregressive frames generation | Self Forcing



SF fix train inference mismatch by conditioning on its own prediction during training.

Self Forcing | Training algorithm

Algorithm 1 Self Forcing Training

Require: Denoise timesteps $\{t_1, \dots, t_T\}$

Require: Number of video frames N

Require: AR diffusion model G_θ (returns KV embeddings via G_θ^{KV})


```
1: loop
2:   Initialize model output  $\mathbf{X}_\theta \leftarrow []$ 
3:   Initialize KV cache  $\mathbf{KV} \leftarrow []$ 
4:   Sample  $s \sim \text{Uniform}(1, 2, \dots, T)$ 
5:   for  $i = 1, \dots, N$  do
6:     Initialize  $x_{t_T}^i \sim \mathcal{N}(0, I)$ 
7:     for  $j = T, \dots, s$  do
8:       if  $j = s$  then
9:         Enable gradient computation
10:        Set  $\hat{x}_0^i \leftarrow G_\theta(x_{t_j}^i; t_j, \mathbf{KV})$ 
11:         $\mathbf{X}_\theta.append(\hat{x}_0^i)$ 
12:        Disable gradient computation
13:        Cache  $\mathbf{kv}^i \leftarrow G_\theta^{\text{KV}}(\hat{x}_0^i; 0, \mathbf{KV})$ 
14:         $\mathbf{KV}.append(\mathbf{kv}^i)$ 
15:      else
16:        Disable gradient computation
17:        Set  $\hat{x}_0^i \leftarrow G_\theta(x_{t_j}^i; t_j, \mathbf{KV})$ 
18:        Sample  $\epsilon \sim \mathcal{N}(0, I)$ 
19:        Set  $x_{t_{j-1}}^i \leftarrow \Psi(\hat{x}_0^i, \epsilon, t_{j-1})$ 
20:      end if
21:    end for
22:  end for
23:  Update  $\theta$  via distribution matching loss
24: end loop
```

Because the student model is causal, KV cache is available, allowing to reuse keys and values in attention.

After generating the output, we run one extra forward at $t=0$ to refresh the KV cache.

Self Forcing | Long video generation via rolling KV cache

conditioning on ALL previous frames (keep full KV cache)



conditioning only on L last frames (evict oldest KV cache)

